GROUNDING AI EXPLANATIONS IN EXPERIENCE: A REFLECTIVE COGNITIVE ARCHITECTURE FOR CLINICAL DECISION SUPPORT

Zijian Shao 2,† , Haiyang Shen 1,3,† , Mugeng Liu 3 , Gecheng Fu 4 , Yaoqi Guo 2 , Yanfeng Wang $^{5,\boxtimes}$, Yun Ma $^{1,3,\boxtimes}$

National Cancer Center/National Clinical Research Center for Cancer/Cancer Hospital, Chinese Academy of Medical Sciences and Peking Union Medical College Beijing, China szj_ngu@stu.pku.edu.cn, hyshen@stu.pku.edu.cn, wangyf@cicams.ac.cn, mayun@pku.edu.cn,

ABSTRACT

Effective disease prediction in modern healthcare demands the twin goals of high accuracy and transparent, clinically meaningful explanations. Existing machine learning and large language model (LLM) based approaches often struggle to balance these goals. Many models yield accurate but unclear statistical outputs, while others generate fluent but statistically unsupported narratives, often undermining both the validity of the explanation and the predictive accuracy itself. This shortcoming comes from a shallow interaction with the data, preventing the development of a deep, detailed understanding similar to a human expert's. We argue that high accuracy and high-quality explanations are not separate objectives but are mutually reinforcing outcomes of a model that develops a deep, direct understanding of the data. To achieve this, we propose the Reflective Cognitive Architecture (RCA), a novel framework that coordinates multiple LLMs to learn from direct experience. RCA features an iterative rule refinement mechanism that improves its logic from prediction errors and a distribution-aware rules check mechanism that bases its reasoning in the dataset's global statistics. By using predictive accuracy as a signal to drive deeper comprehension, RCA builds a strong internal model of the data. We evaluated RCA on one private and two public datasets against 22 baselines. The results demonstrate that RCA not only achieves state-of-the-art accuracy and robustness with a relative improvement of up to 40% over the baseline but, more importantly, leverages this deep understanding to excel in generating explanations that are clear, logical, evidence-based, and balanced, highlighting its potential for creating genuinely trustworthy clinical decision support systems. The code is available at https://github.com/ssssszj/RCA. Access to the CRT dataset requires an application and approval process. If approved, we will anonymize and open-source the data.

1 Introduction

Disease prediction is a foundation of modern healthcare, providing a crucial opportunity for timely interventions that can slow disease progression, improve patient outcomes, and reduce medical costs (Nahian et al., 2022). In clinical practice, the data for such predictions are typically structured in tabular formats, containing a wealth of patient information (Nahian et al., 2022; Fang et al., 2024). However, the ultimate usefulness of a predictive model in a high-stakes clinical setting is determined

¹Institute for Artificial Intelligence, Peking University

²School of Software & Microelectronics, Peking University

³School of Computer Science, Peking University

⁴ School of Life Sciences, Peking University

⁵Department of Comprehensive Oncology,

[†]Equal Contribution, [®]Corresponding Author

by the twin requirements of predictive accuracy and the ability to explain its reasoning in a manner that clinicians can trust and act upon.

The major gap in current predictive systems lies in meeting both these needs at the same time. An effective system must not only predict correctly but also generate high-quality descriptions that explain its reasoning. Such an explanation must meet several key criteria derived from cognitive science and medical practice. First, it must have a low Cognitive Load (CL) (Sweller, 2011), presenting information clearly and concisely. Second, it must show sound Logical Argumentation (LA) (Toulmin, 2003), with a coherent reasoning process. Third, it must be based on Evidence-based Medicine (EBM) (Guyatt et al., 1992), matching both established medical knowledge and the statistical facts of the data. Finally, it must actively reduce Cognitive Biasing (CB) (Kahneman, 2011) by presenting a balanced view. The failure of current AI to deliver accurate predictions paired with such high-quality explanations is a main obstacle to its adoption in clinical decision-making.

The ongoing challenge is that existing methods often fail to meet both requirements. Classical machine learning models, such as linear regression (Tibshirani, 1996) and tree-based approaches (Prokhorenkova et al., 2018), can achieve good results, but their explanatory ability is limited to statistical outputs like feature importance scores. These are not narrative explanations and need significant expert analysis, increasing cognitive load. On the other hand, the arrival of Large Language Models (LLMs) (Zhao et al., 2025) brought the promise of natural language explanations. However, when applied simply, they often lack a deep, detailed understanding of the specific dataset. Their reasoning can become "statistically unsupported," a weakness that leads to two problems: they produce explanations that seem medically believable but are not supported by the data, and this same shallow understanding often harms their predictive accuracy.

Our key insight is that high predictive accuracy and high-quality explanations are not conflicting goals but are two results of a single, deeper process: developing a direct, experience-based understanding of the data. This is like how a human expert dives deep into data before drawing conclusions. We rethink predictive accuracy not just as an end goal, but as a crucial reward signal that drives the model to build a more robust and fundamental "experience" with the underlying patterns. By optimizing for correct predictions and improving robustness against the data noise common in medical datasets, the model is forced to achieve a deep data understanding. This deep understanding is the necessary condition for generating explanations that are insightful, reliable, and clinically useful, with a high-performance predictive model appearing as a valuable, simultaneous output.

To achieve this, we propose the **Reflective Cognitive Architecture (RCA)**, a framework designed to enable LLMs to learn directly from data through a process of experience and reflection. RCA includes two core mechanisms. The iterative rule refinement mechanism allows the model to learn from its mistakes, treating each incorrect prediction as an experience to be turned into abstract rules, thereby creating sound logical argumentation (LA). In addition to this, the distribution-aware rules check mechanism makes sure these rules are statistically based in the overall data distribution, promoting evidence-based reasoning (EBM) and reducing cognitive biases (CB).

We conducted extensive experiments on three disease prediction datasets, including a private real-world dataset for Catheter-Related Thrombosis (CRT), comparing RCA against 22 baselines. Our evaluation judges models on their ability to achieve both high predictive performance and high-quality explanations, as well as their robustness to data noise. The results demonstrate that RCA significantly outperforms existing methods on all fronts, confirming our main idea that a deeper, experience-based understanding of data is the key to achieving truly explainable and accurate AI for healthcare.

In summary, our main contributions are as follows:

- We rethink the problem of explainable disease prediction, arguing that predictive accuracy and high-quality descriptive statements are mutually reinforcing outcomes of a deep data understanding, which should be the main goal of the model.
- We propose RCA, A novel architecture featuring an iterative refinement mechanism to build understanding from experience and a distribution-aware check mechanism to make sure this understanding is statistically based and robust.
- We prepare a real-world dataset for CRT and create a complete evaluation framework judging
 predictive accuracy, robustness, and explanation quality based on principles of cognitive load,
 logical argumentation, evidence-based medicine, and cognitive bias.
- We demonstrate through extensive experiments that RCA outperforms a wide range of baselines in achieving a better balance of accuracy and explanation quality, supported by good robustness.

2 RELATED WORK

In this section, we place our work within two key areas: the evolution of explainability in disease prediction models (2.1) and the data interaction methods of LLM-based agents (2.2).

2.1 EXPLAINABILITY IN DISEASE PREDICTION

The quest for explainability in disease prediction is not new, but its definition and relationship with accuracy have changed over time (Sun et al., 2024). Early approaches favored models that were naturally understandable. For instance, methods like linear regression (Hoerl & Kennard, 2000; Tibshirani, 1996) and decision trees (Breiman, 2001; Prokhorenkova et al., 2018) were valued for their transparency and provided reasonable performance. However, their explanations were limited to statistical outputs, not narrative descriptions, which require significant expert analysis and place a high cognitive load on physicians.

Subsequently, more complex deep learning models emerged, achieving very high accuracy using Transformer (Hollmann et al., 2022) or FFN (Gorishniy et al., 2023). But most of them were often "black boxes," making their reasoning difficult to understand. This shifted the focus to post-hoc explanation techniques. However, these can be unfaithful to the model's actual reasoning process. The emergence of LLMs offered a path toward generating natural language explanations. Initial efforts used LLMs to interpret statistical outputs or directly analyze tabular data. While these methods can produce fluent text, they often suffer from a disconnect from the data's statistics. This lack of grounding is a critical weakness, as it frequently leads to a dual failure: the explanations are not only invalid and untrustworthy, but the shallow reasoning process also harms the accuracy of the prediction itself. Our work addresses this by creating a direct link between the learning process for accuracy and the generation of explanations, ensuring they are two sides of the same coin.

2.2 LLM-BASED AGENT

LLM-based agents are increasingly being developed to tackle complex tasks by interacting with external environments or tools (Shen, 2024; Wang et al., 2024; Wu et al., 2024). The main methods for data analysis tasks involve optimizing for tool use or code generation.

API-based agents (Shen, 2024; SHEN et al., 2025; Shen et al., 2025) interact with data through a fixed set of functions. This approach creates a layer of abstraction between the agent and the raw data. The agent learns to become a skilled tool-caller, but it does not develop a detailed, instance-level understanding of the data's specifics. This abstraction hinders its ability to generate descriptive statements that are rich in detail, as it only perceives the data through the summarized lens of its tools.

Code-generation agents (OpenAI, 2023; Zhang et al., 2024; Guo et al., 2024) represent a more flexible approach, where the LLM writes and executes code (e.g., Python scripts) to perform analysis. While powerful, this method still maintains a degree of separation. The agent's core task becomes generating correct code, and the insights are derived from the code's output. The process of deep, repeated reflection on individual data points and their relationships is often bypassed in favor of executing a script that provides a summarized result.

RCA deliberately departs from these tool-focused methods. It is an agent that engages with data directly, without the mediation of external tools or code interpreters. Its core innovation lies in its internal, reflective process, which forces the model to build its understanding from direct "experience" with the data, imitating how a human analyst develops intuition. By avoiding the shortcuts of tool use, RCA optimizes for data comprehension directly, fostering a deeper and more robust understanding that serves as the essential foundation for both high accuracy and high-quality explanations.

3 Methods

Our methodology is designed to build a model that is both highly accurate and produces superior explanatory statements. We posit that this dual objective is best achieved by forcing the model to develop a deep, experience-driven understanding of the data. To this end, we designed the Reflective Cognitive Architecture (RCA), a framework that fosters this deep engagement rather than a superficial

analysis. The final output for each patient s_i is a prediction $\hat{\mathcal{Y}}_i$ that includes not only an accurate binary disease label $\hat{y}_i \in \{0,1\}$ but also a high-quality explanatory statement \hat{e}_i . Formally, let $S = \{s_i\}_{i=1}^N$ be the dataset where each patient is represented by a vector of structured clinical features f_i and a true disease state y_i . This section first provides an overview of the RCA architecture and then details its two core components: the Iterative Rules Optimization mechanism and the Distribution-aware Rules Check mechanism.

3.1 RCA OVERVIEW

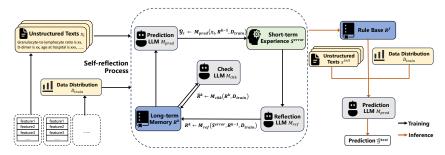


Figure 1: RCA pipeline. RCA uses reflective cycles and additional checks to directly analyze data, building a deep understanding that enhances both prediction accuracy and the generation of detailed, grounded explanatory statements.

The RCA pipeline, shown in Figure 1, is engineered to systematically build and refine a deep data understanding, which is the necessary precondition for achieving high accuracy and generating high-quality explanations. The process unfolds as follows:

Data Narrative. To enable an LLM to "read" and reason about patient data, we first transform the structured features f_i into an unstructured text narrative x_i as shown in Figure 1. For example, a data row is converted into a sentence like "Granulocyte-to-lymphocyte ratio is 4.88, D-dimer is 3.16, no chemotherapy, catheterization is CVC." This makes the data directly accessible to the LLM, with the patient record reformulated as $s_i = (x_i, y_i)$.

Distribution Extraction. To ensure the model's reasoning is grounded in Evidence-based Medicine (EBM) and to mitigate Cognitive Biasing (CB), we provide it with a global statistical context. Along with data transformation in Figure 1, we extract the data distribution \mathcal{D}_{train} from the training set, which summarizes the statistical properties of the entire patient cohort (e.g., means, quantiles, frequencies). This global context prevents the model from over-interpreting individual data points.

Guided Prediction via a Dynamic Rule Base. At the heart of RCA is a dynamic rule base, R, which serves as the model's evolving long-term memory. The "Self-reflection Process" of Figure 1 shows that, for any given patient x_i , a prediction LLM, M_{pred} , uses the current rule base R^{k-1} and the global data distribution \mathcal{D}_{train} to generate its output:

$$\hat{\mathcal{Y}}_i = (\hat{y}_i, \hat{e}_i) \leftarrow M_{pred}(x_i, R^{k-1}, \mathcal{D}_{train})$$
(1)

By explicitly using the rule base, the model's predictions are guided by a consistent set of principles, fostering sound Logical Argumentation (LA) within the explanation \hat{e}_i .

Training and Prediction Phases. The training phase is a dynamic process where the rule base is iteratively refined by two other LLMs, M_{ref} and M_{chk} , based on prediction outcomes. In the testing phase, marked with orange lines in Figure 1, the final, optimized rule base R^f is used to generate predictions and explanations $\hat{\mathcal{Y}}^{test}$ for unseen data. The resulting explanation is thus a product of a deep, iterative learning process that was driven by the pursuit of accuracy.

3.2 ITERATIVE RULES OPTIMIZATION: BUILDING LOGICAL ARGUMENTATION FROM EXPERIENCE

To generate explanations with sound Logical Argumentation (LA), a model must possess a coherent reasoning framework. The iterative rules optimization mechanism builds this framework by emulating

experiential learning: treating prediction errors as opportunities to reflect and refine its understanding. This mechanism converts the short-term experience of specific errors into the long-term memory of abstract, generalizable rules.

Iterative Reflection Loop. The process is a feedback loop where M_{pred} uses the rule base R^{k-1} to make predictions. The instances where its predictions are incorrect constitute the direct feedback—the "experience"—that drives learning.

Short-term Experience via Error Samples. Misclassified samples are collected into a textual format, S^{error} :

$$S^{error} = \operatorname{conc}(s_j^{error})_{j=1}^T \tag{2}$$

where T is the error batch capacity. This aggregation of incorrect cases highlights deficiencies in the current rule base.

Long-term Memory in the Rule Base. This experience is processed by a reflection LLM, M_{ref} , to update the rule base, distilling specific errors into robust principles:

$$R^k \leftarrow M_{ref}(S^{error}, R^{k-1}, \mathcal{D}_{train}) \tag{3}$$

Through this process, the model's logical framework (R^k) evolves through trial and error. This ensures that the pursuit of higher accuracy directly sharpens the logical rules that will form the backbone of the final explanation.

3.3 DISTRIBUTION-AWARE RULES CHECK: GROUNDING LOGIC IN EVIDENCE

While iterative learning builds a logical framework, it risks creating rules that are statistically spurious. To ensure explanations are grounded in Evidence-based Medicine (EBM), mitigate Cognitive Biasing (CB), and improve robustness, the logic must be validated against the global data. The distribution-aware rules check mechanism serves as a safeguard, ensuring the model's reasoning is statistically sound.

Additional Rules Check. At the end of each epoch, a checking LLM, M_{chk} , reviews the rule base R^k using the global data distribution \mathcal{D}_{train} as a reference:

$$\hat{R}^k \leftarrow M_{chk}(R^k, \mathcal{D}_{train}) \tag{4}$$

 M_{chk} removes low-quality or overly specific rules and summarizes general rules for detecting outliers. This grounds the model's reasoning in the statistical properties of the dataset, directly promoting an evidence-based approach and strengthening the model against noisy or atypical data. The refined rule base \hat{R}^k then replaces the previous version for the next epoch (denoted as R^k for consistency).

Mutual Enhancement. Together, the iterative optimization and distribution-aware check form a synergistic, closed-loop system:

$$\hat{R}^{k} \overset{cover}{\underset{M_{chk}}{\rightleftharpoons}} R^{k} \overset{M_{pred}}{\underset{M_{ref}}{\rightleftharpoons}} \hat{\mathcal{Y}}_{i} \tag{5}$$

The iterative process (M_{pred}, M_{ref}) builds the core logical structure (LA) from the experience of pursuing accuracy, while the check mechanism (M_{chk}) ensures this structure is statistically robust and evidence-based (EBM, CB). This dual-process architecture ensures that RCA develops a deep, reliable understanding of the data, which is the essential foundation for achieving both high accuracy and generating trustworthy explanatory statements.

4 EVALUATION

We designed our evaluation to test the central hypothesis: that a deeper, experience-driven data understanding leads to synergistic improvements in predictive accuracy, explanation quality, and robustness. For clinical decision support, both accuracy and explanations are critical. This section first details the experimental setup (4.1), then presents the main results correlating performance and explanation quality (4.2), tests the model's resilience against data noise (4.3), validates our architecture through an ablation study (4.4), and provides a qualitative case study (4.5).

4.1 SETUP

Datasets. To ensure a comprehensive evaluation, we selected three distinct datasets. For each dataset, we split it into training, validation, and test sets following a 3:1:1 ratio.

- **CRT:** We curated a real-world dataset for Catheter-Related Thrombosis (CRT) in collaboration with Feitian Hospital¹. This proprietary dataset comprises 315 cancer patients, offering a high-stakes, clinically relevant challenge. This research was approved by the Medical Science Research Ethics Committee of the authors' institute.
- **Diabetes** (Pore, 2025): A public benchmark dataset for diabetes prediction with 8 highly correlated features. We use a subset of 415 cases to test the model's performance on a well-understood clinical problem.
- **Heart Disease** (Rdeki, 2025): A public dataset for heart disease prediction featuring 19 primarily categorical features, including lifestyle and biometric data. Its 965 cases challenge the model's ability to reason over heterogeneous data types.

More details of datasets are provided in Appendix A.4.

Baselines. We compare RCA against 22 baselines representing diverse approaches.

- Traditional ML Models: We include Lasso regression (Tibshirani, 1996) and Catboost (Prokhorenkova et al., 2018). These models are standard for tabular data but produce statistical artifacts (e.g., coefficients, feature importance) rather than narrative explanations, representing a baseline for expert-driven interpretation. A 'Qwen3-235B' is used to generate an explanation for these models; we provide a detailed introduction in Appendix A.5.
- LLM-based Methods: We test the ability of 4 leading non-reasoning LLMs ('Qwen2.5-72B-Instruct', 'DeepSeek-V3-64k', 'DeepSeek-Chat-V3.1', 'GPT-4.1-2025-04-14') to perform zero-shot prediction and explanation directly from tabular data.
- **Reasoning LLMs:** We evaluate 6 advanced reasoning-focused LLMs ('DeepSeek-R1', 'Qwen3-30B-A3B', 'Qwen3-235B-A22B-Instruct-2507', 'GPT-5-2025-08-07', 'o3-mini-2025-01-31', 'o4-mini-2025-04-16') to assess whether enhanced general reasoning capabilities translate to better data understanding and explanation in this specific domain.
- **LLM-based Agents:** We include two agent paradigms that reflect the state-of-the-art in LLM-driven analysis, comprising 9 and 1 methods, respectively. The **LLM+Tools** approach equips models with predefined functions, while the **LLM+Code** approach utilizes a code interpreter (OpenAI, 2023). These baselines test the hypothesis that tool use abstracts away the fine-grained data interaction necessary for deep understanding. For details about the functions and code, please refer to the Appendix A.6.

Evaluation Metrics. Our evaluation employs two categories of metrics.

- Metrics for Predictive Performance: We use accuracy, the Matthews Correlation Coefficient (MCC), and the F1-score. These metrics serve as quantitative proxies for the depth and correctness of the model's understanding, with MCC and F1-score being particularly informative for the imbalanced datasets common in medicine.
- Metrics for Explanation Quality: To directly measure the quality of the generated descriptive statements, we developed four criteria grounded in cognitive science and medical practice: Cognitive Load (CL), Logical Argumentation (LA), Evidence-based Medicine (EBM), and Cognitive Biasing (CB). These criteria have been recognized by 3 doctors as clinically valuable, as they align with real-world medical assessment needs for evaluating the clarity, rationality, and reliability of explanatory content. Then doctors are invited to score each explanation on a scale of 1 to 10 for each criterion based on a detailed rubric. The rubric of these metrics are provided in Appendix A.7, while specific examples can be found in Appendix A.8.

Implementation Details. For our method, we implement RCA using both 'Qwen2.5-72B-Instruct' and 'GPT-4.1-2025-04-14' as the base LLMs(abbreviated as RCA+Qwen2.5 and RCA+GPT-4.1 to

¹The hospital name has been anonymized to comply with the anonymity policy.

demonstrate its architectural benefits. The error batch capacity T is set to 25. The model is trained for 15 epochs on the CRT and Diabetes datasets, and 25 epochs on the Heart Disease dataset. All prompt templates used in RCAcan be found in Appendix A.12

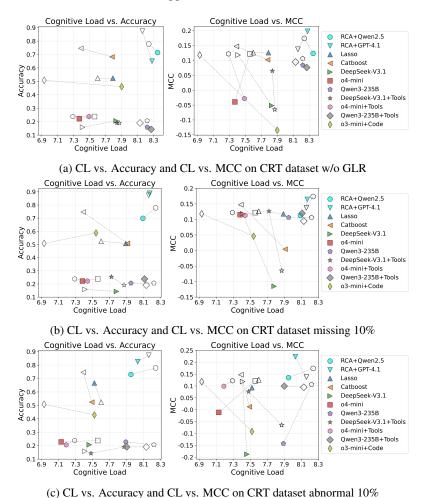


Figure 2: Results of the main experiment and the robustness experiment on CRT dataset. Hollow dots represent the main experiment, solid dots represent the robustness experiment, and dots of the same shape represent the same approach. The dashed line measures performance variation. RCA demonstrates not only the best results (4.2) but also gets little performance fluctuations(4.3, showing the resilience to data noise.

4.2 MAIN RESULTS: DEEP UNDERSTANDING YIELDS SYNERGISTIC GAINS

In this section, we conduct qualitative analysis of main experiment on the CRT dataset, while the quantitative results of all datasets can be found in the Appendix A.9. Specifically, the qualitative analysis includes Lasso, Catboost, 'DeepSeek-Chat-V3.1'(abbreviated as 'DeepSeek-V3.1'), 'o4-mini-2025-04-16'(abbreviated as 'o4-mini'), 'Qwen3-235B-A22B-Instruct-2507' (abbreviated as 'Qwen3-235B') and 'o3-mini-2025-01-31+Code' (abbreviated as 'o3-mini+Code') as baselines.

Our central thesis posits that predictive accuracy and explanation quality are not in opposition but are synergistic outcomes of a model's deep, first-hand understanding of data. We test this by plotting model performance against explanation quality. As shown by the hollow points in Figure 2, which represent the main experimental results, the RCA-based approaches consistently occupy the top-right quadrant, achieving the highest predictive performance (Accuracy and MCC) while also delivering high-quality explanations (as indicated by the Cognitive Load score, with detailed metrics in Appendix A.9).

This result validates our core hypothesis. RCA's success is not coincidental; it is a direct consequence of its architecture, which forces a deep engagement with the data. The high predictive accuracy serves as confirmation that the internal model RCA has built is a correct representation of the underlying clinical patterns. The high-quality explanation is the natural articulation of this well-grounded understanding.

In contrast, other paradigms falter because they lack this deep engagement.

- Traditional ML models like Catboost achieve competitive accuracy, but explanations generated by 'Qwen3-235B' based on their prediction results perform poorly overall.
- LLM-based agents (LLM+Tools, LLM+Code) are hindered by their layer of abstraction. By interacting with data through APIs or code outputs, they become proficient tool-users but never develop a granular, instance-level "feel" for the data. Their understanding remains second-hand, limiting the depth and reliability of their explanations.
- Standalone LLMs, even powerful reasoning models like 'o4-mini', demonstrate the pitfalls of "statistically de-grounded" reasoning. While they may achieve a respectable MCC, their tendency to generate plausible but unsubstantiated narratives leads to lower accuracy and poorer explanation scores, highlighting a superficial grasp of the specific dataset.

Ultimately, the results show that by prioritizing the development of a deep data understanding, RCA organically achieves state-of-the-art results in both prediction and explanation, demonstrating their synergistic relationship. Full results for all baselines across three datasets please refer to Appendix A.9.

4.3 ROBUSTNESS: DEEP UNDERSTANDING CONFERS RESILIENCE

Similarly, we conduct qualitative analysis of robust experiment on the CRT dataset, while the quantitative results can be found in the Appendix A.10.

Medical data is notoriously noisy and incomplete. A truly deep understanding should be resilient to such imperfections, distinguishing robust signals from spurious noise. We tested this by degrading the CRT dataset through feature removal ("GLR"), random value deletion (10%), and the introduction of outliers (10%).

The results, visualized by the solid points and connecting dashed lines in Figure 2, demonstrate RCA's superior robustness. The performance of RCA (both 'Qwen2.5-72B' and 'GPT-4.1' versions) shows minimal degradation across all noise conditions, as indicated by the short dashed lines connecting the hollow (original) and solid (noisy) points. This stability is a direct outcome of its design. The distribution-aware rules check grounds the model's logic in global statistics, preventing it from being misled by local anomalies or outliers. The iterative rules optimization builds a generalized understanding from cumulative experience, making the model less dependent on any single data point or feature.

In contrast, the performance of many baselines is far more volatile. For instance, both Catboost and 'DeepSeek-V3.1' suffer significant drops in MCC when faced with missing data, revealing that their underlying models may have overfit to patterns that are not robust. Their longer dashed lines signify a shallower understanding that shatters under data stress. The resilience of RCA is therefore not just a feature but further evidence of the foundational and robust nature of its data understanding.

4.4 ABLATION STUDY

To validate that RCA's performance stems directly from its proposed cognitive mechanisms, we conducted an ablation study by systematically removing its core components. As shown in Table 1, the removal of any key module leads to a significant performance collapse, confirming that each part is essential to the process of building a deep understanding. More explanation results are provided in Appendix A.11.

Table 1: Results of ablation studies. The experimental results indicate that several core modules in RCA play irreplaceable roles.

		Qwen2	.5-72B			GPT	-4.1					
	original	distribution	reflection	check	ck original distribution reflection		check					
	CRT											
Accuracy MCC F1-score CL	0.7778 0.1739 0.2222 8.24	0.5873 -0.0702 0.0714 7.70	0.5714 0.1513 0.1818 7.54	0.6032 0.1691 0.1935 7.75	0.8730 0.1373 0.2000 8.16	0.6508 0.0824 0.1538 7.45	0.7143 -0.0024 0.1000 7.57	0.7937 0.0517 0.1333 7.40				
Diabetes												
Accuracy MCC F1-score CL	0.7831 0.5406 0.7097 8.13	0.7711 0.4926 0.6667 7.93	0.7229 0.4424 0.6567 7.93	0.7349 0.4169 0.6206 7.77	0.7470 0.4244 0.6038 8.03	0.7229 0.3857 0.5965 7.18	0.6867 0.3222 0.5667 7.34	0.7349 0.4232 0.5652 7.73				
			He	art Diseas	se							
Accuracy MCC F1-score CL	0.5647 0.0547 0.1290 7.62	0.3523 -0.0477 0.3169 7.32	$0.3575 \\ -0.1001 \\ 0.2874 \\ 7.54$	0.3627 -0.0349 0.3204 7.53	0.7461 0.1493 0.2898 7.74	0.5337 -0.0359 0.2623 7.12	0.3471 -0.1134 0.2841 7.13	0.4663 -0.0923 0.2481 6.72				

4.5 QUALITATIVE ANALYSIS: A PREDICTION CASE STUDY

Aggregate metrics validate our approach, but a case study reveals the practical difference between superficial and deep understanding. Figure 3 contrasts the explanations from RCA and a strong reasoning baseline, 'DeepSeek-R1', for the same patient from the CRT dataset.

'DeepSeek-R1' incorrectly predicts CRT, exemplifying the danger of statistically de-grounded reasoning. It constructs a plausible-sounding narrative by identifying risk factors (CVC, chemotherapy) but critically fails in its quantitative assessment. It misinterprets a D-dimer level of 0.84 mg/L as high risk, demonstrating a lack of awareness of the actual risk thresholds learned from the data distribution. This is a classic failure of a model that relies on general knowledge rather than a first-hand understanding of the specific dataset's statistical realities.

In contrast, RCA correctly predicts no CRT and generates an explanation that is a direct manifestation of its deep data understanding.

- Its reasoning is grounded in **Evidence-based Medicine (EBM)**, a result of the *distribution-aware rules check*. It explicitly compares the patient's D-dimer (0.84) and GLR (0.5) against the learned clinical risk thresholds (e.g., >1.5), correctly concluding they are "notably lower" and "well below the risk threshold."
- It exhibits strong **Logical Argumentation** (**LA**), a product of the *iterative rules optimization*. It presents a balanced view, acknowledging risk factors but correctly reasoning that they are "insufficient for thrombosis without an elevated D-dimer."

This clear, logical, and evidence-based explanation is not a separate feature but the output of the same deep understanding that drove the accurate prediction. It is precisely this synergy that makes RCA a step towards truly trustworthy clinical AI.

5 Conclusion

In this paper, we challenged the conventional trade-off between predictive accuracy and explainability in clinical AI. We argued that these are not competing goals but synergistic outcomes of a model that develops a deep, first-hand understanding of the data. To achieve this, we introduced RCA, a novel framework that learns from experience through iterative rules optimization and grounds its reasoning in global statistics via a distribution-aware check. Our experiments demonstrate that by forcing the model to achieve this deeper comprehension, RCA not only attains state-of-the-art accuracy and robustness but also excels in generating clear, logical, and evidence-based explanatory statements.

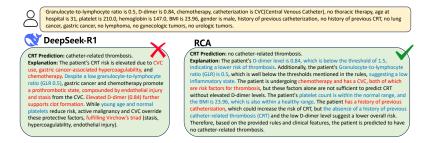


Figure 3: Comparison of explanations from 'DeepSeek-R1' and RCA for the same patient. RCA demonstrates superior reasoning by integrating quantitative thresholds and providing a balanced, evidence-based argument, a direct result of its deep data understanding. 'DeepSeek-R1' sexplanation, while fluent, is statistically ungrounded and leads to an incorrect prediction.

REFERENCES

- Leo Breiman. Random forests. *Mach. Learn.*, 45(1):5–32, October 2001. ISSN 0885-6125. doi: 10.1023/A:1010933404324. URL https://doi.org/10.1023/A:1010933404324.
- Xi Fang, Weijie Xu, Fiona Anting Tan, Ziqing Hu, Jiani Zhang, Yanjun Qi, Srinivasan H. Sengamedu, and Christos Faloutsos. Large language models (LLMs) on tabular data: Prediction, generation, and understanding a survey. *Transactions on Machine Learning Research*, 2024. ISSN 2835-8856. URL https://openreview.net/forum?id=IZnrCGF9WI.
- Yury Gorishniy, Ivan Rubachev, Nikolay Kartashev, Daniil Shlenskii, Akim Kotelnikov, and Artem Babenko. Tabr: unlocking the power of retrieval-augmented tabular deep learning. *CoRR*, 2023.
- Chengquan Guo, Xun Liu, Chulin Xie, Andy Zhou, Yi Zeng, Zinan Lin, Dawn Song, and Bo Li. Redcode: Risky code execution and generation benchmark for code agents. *Advances in Neural Information Processing Systems*, 37:106190–106236, 2024.
- Gordon Guyatt, John Cairns, David Churchill, Deborah Cook, Brian Haynes, Jack Hirsh, Jan Irvine, Mark Levine, Mitchell Levine, Jim Nishikawa, et al. Evidence-based medicine: a new approach to teaching the practice of medicine. *jama*, 268(17):2420–2425, 1992.
- Arthur E. Hoerl and Robert W. Kennard. Ridge regression: biased estimation for nonorthogonal problems. *Technometrics*, 42(1):80–86, February 2000. ISSN 0040-1706. doi: 10.2307/1271436. URL https://doi.org/10.2307/1271436.
- Noah Hollmann, Samuel Müller, Katharina Eggensperger, and Frank Hutter. Tabpfn: A transformer that solves small tabular classification problems in a second. *arXiv preprint arXiv:2207.01848*, 2022.
- Daniel Kahneman. Thinking, fast and slow. Farrar, Straus and Giroux, 2011.
- Jabir Al Nahian, Abu Kaisar Mohammad Masum, Sheikh Abujar, and Md. Jueal Mia. Common human diseases prediction using machine learning based on survey data, 2022. URL https://arxiv.org/abs/2209.10750.
- OpenAI. Data analysis with chatgpt, 2023. URL https://help.openai.com/en/articles/8437071-data-analysis-with-chatgpt.
- Nandita Pore. Healthcare diabetes dataset. https://www.kaggle.com/datasets/nanditapore/healthcare-diabetes/data, 2025. Accessed: 2025-05-15.
- Liudmila Prokhorenkova, Gleb Gusev, Aleksandr Vorobev, Anna Veronika Dorogush, and Andrey Gulin. Catboost: unbiased boosting with categorical features. In *Proceedings of the 32nd International Conference on Neural Information Processing Systems*, NIPS'18, pp. 6639–6649, Red Hook, NY, USA, 2018. Curran Associates Inc.
- Oktay Rdeki. Heart disease dataset, 2025. URL https://www.kaggle.com/datasets/oktayrdeki/heart-disease. Accessed: 2025-05-15.

- Haiyang SHEN, Yue Li, Desong Meng, Dongqi Cai, Sheng Qi, Li Zhang, Mengwei Xu, and Yun Ma. Shortcutsbench: A large-scale real-world benchmark for API-based agents. In *The Thirteenth International Conference on Learning Representations*, 2025. URL https://openreview.net/forum?id=kKILfPkhSz.
- Haiyang Shen, Hang Yan, Zhongshi Xing, Mugeng Liu, Yue Li, Zhiyang Chen, Yuxiang Wang, Jiuzheng Wang, and Yun Ma. Ragsynth: Synthetic data for robust and faithful rag component optimization, 2025. URL https://arxiv.org/abs/2505.10989.
- Zhuocheng Shen. Llm with tools: A survey, 2024. URL https://arxiv.org/abs/2409.18807.
- Qiyang Sun, Alican Akman, and Björn W. Schuller. Explainable artificial intelligence for medical applications: A review, 2024. URL https://arxiv.org/abs/2412.01829.
- John Sweller. Chapter two cognitive load theory. In Jose P. Mestre and Brian H. Ross (eds.), *Psychology of Learning and Motivation*, volume 55 of *Psychology of Learning and Motivation*, pp. 37–76. Academic Press, 2011. doi: https://doi.org/10.1016/B978-0-12-387691-1.00002-8. URL https://www.sciencedirect.com/science/article/pii/B9780123876911000028.
- Robert Tibshirani. Regression shrinkage and selection via the lasso. *Journal of the Royal Statistical Society. Series B (Methodological)*, 58(1):267–288, 1996. ISSN 00359246. URL http://www.jstor.org/stable/2346178.
- Stephen E Toulmin. The uses of argument. Cambridge university press, 2003.
- Jize Wang, Ma Zerun, Yining Li, Songyang Zhang, Cailian Chen, Kai Chen, and Xinyi Le. Gta: a benchmark for general tool agents. In *The Thirty-eight Conference on Neural Information Processing Systems Datasets and Benchmarks Track*, 2024.
- Shirley Wu, Shiyu Zhao, Qian Huang, Kexin Huang, Michihiro Yasunaga, Kaidi Cao, Vassilis Ioannidis, Karthik Subbian, Jure Leskovec, and James Y Zou. Avatar: Optimizing Ilm agents for tool usage via contrastive reasoning. *Advances in Neural Information Processing Systems*, 37: 25981–26010, 2024.
- Kechi Zhang, Jia Li, Ge Li, Xianjie Shi, and Zhi Jin. CodeAgent: Enhancing code generation with tool-integrated agent systems for real-world repo-level coding challenges. In Lun-Wei Ku, Andre Martins, and Vivek Srikumar (eds.), *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 13643–13658, Bangkok, Thailand, August 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.acl-long.737. URL https://aclanthology.org/2024.acl-long.737/.
- Wayne Xin Zhao, Kun Zhou, Junyi Li, Tianyi Tang, Xiaolei Wang, Yupeng Hou, Yingqian Min, Beichen Zhang, Junjie Zhang, Zican Dong, Yifan Du, Chen Yang, Yushuo Chen, Zhipeng Chen, Jinhao Jiang, Ruiyang Ren, Yifan Li, Xinyu Tang, Zikang Liu, Peiyu Liu, Jian-Yun Nie, and Ji-Rong Wen. A survey of large language models, 2025. URL https://arxiv.org/abs/2303.18223.

A APPENDIX / SUPPLEMENTAL MATERIAL

A.1 DECLARATION ON THE USE OF LLMS

We declare that the use of LLMs during the preparation of this manuscript was strictly limited to language-related assistance, such as sentence refinement and grammatical correction. All substantive content was independently authored by the authors and rigorously reviewed and verified following any LLM-assisted modifications. During the experiments, all usage of LLMs was solely for academic research purposes, with no inappropriate applications. Detailed experimental settings are provided in the Experiments section of this paper. No other reliance on LLMs is involved in this work.

A.2 ETHICS STATEMENT

All authors of this paper strictly adhere to the ICLR Code of Ethics throughout the entire research and manuscript preparation process. Prior to submitting this work, every author has carefully read and fully understood the content of the Code of Ethics, and confirms compliance with all its provisions.

This research strictly follows the ethical guidelines for conference participation outlined in the ICLR Code of Ethics, covering the integrity of paper submission, adherence to ethical standards in potential peer review , and constructive, respectful communication in any subsequent academic discussions related to this work. We further confirm that no content in this paper violates the ethical principles specified in the ICLR Code of Ethics, and that all research procedures are conducted with honesty, transparency, and respect for academic ethics.

A.3 REPRODUCIBILITY STATEMENT

The work in this paper is well reproducible. The code used in the study is available at https://anonymous.4open.science/r/anonym107. Among the three datasets employed in the experiments, the Diabetes Dataset(https://www.kaggle.com/datasets/nanditapore/healthcare-diabetes/data) and Heart Disease Dataset(https://www.kaggle.com/datasets/oktayrdeki/heart-disease) are public and can be accessed via the corresponding links; only the CRT Dataset is private, and its access requires submission of an application and subsequent approval. Data processing can refer to Section 3.1.

A.4 DATASET DETAILS

In this section, we will provide a detailed overview of the dataset.

- CRT: We collected a real dataset on Catheter-Related Thrombosis (CRT) for cancer patients from Feitian Hospital The dataset includes 315 cancer patients who underwent catheterization, with 32 diagnosed with CRT. The dataset contains a total of 17 features, 11 of which are categorical features and 6 are numerical features, including various tumor labels, laboratory test values, and other medically relevant data.
- Diabetes: This public dataset includes 8 features that are strongly associated with diabetes. We randomly selected 415 cases for diabetes prediction, of which 153 had diabetes. Among these 8 features, only "number of pregnancies" is a categorical feature, while the remaining 7 are numerical features.
- Heart Disease: This is a heart disease prediction dataset consisting of 19 features, most of which
 are categorical features. These features include lifestyle habits, blood tests, etc. We have 965
 cases in this dataset, 193 of which were diagnosed with heart disease. In this dataset, numerical
 features are dominated by various laboratory test values, while categorical features mainly cover
 indicators related to different living habits.

A.5 EXPLANATION OF TRADITIONAL MLS

Unlike LLMs, traditional machine learning algorithms cannot directly produce textual explanations. Therefore, for Lasso Regression and CatBoost, we input the prediction results, along with the feature coefficients or feature importance produced by each method, into a Qwen3-235B, which performs best among all LLMs in explanation scores. Then It's prompted to generate explanations based on prediction results and coefficients or importance.

A.6 LLM-BASED AGENTS

A.6.1 LLM+Tools

We pre-defined a logistic regression function and a decision tree function to conduct relevant analysis,. These two functions can accept sample data and output feature correlation coefficients (for the logistic regression function) or feature importance (for the decision tree function) based on their built-in logic. For the LLM+tools method, when the LLM determines that tool invocation is necessary, it first organizes the data into a pre-specified format and sends it to the tool; after receiving the results returned by the tool, the LLM further uses these results to achieve data interpretation.

A.6.2 LLM+CODE

The code interpreter of o3-mini is one of its core functional modules, supporting dynamic tool generation and code execution. As a cost-effective inference model released by OpenAI, o3-mini exhibits excellent performance in STEM (Science, Technology, Engineering, Mathematics) fields—including science, mathematics, and coding (OpenAI, 2023). To this end, we input data into the LLM + code agent, and explicitly instruct the LLM to analyze the target problem by writing code; this setup aims to investigate whether the code interpreter can enhance the LLM's understanding of the data.

A.7 RUBRIC OF EXPLANATION CRITERIA

Considering the value of clinical application, We designed detailed scoring rubric on each criterion for doctors to reference. Cognitive Load, Logical Argumentation, Evidence-Based Medicine and Cognitive Biasing are presented sequentially in Section

A.7.1 COGNITIVE LOAD(CL)

Considering that the explanation is ultimately read by doctors, the definition of Cognitive Load is "Whether the explanation is easy for doctors to understand and analyze". Specific standards are as follows:

- **7-10 points:** Extremely easy to understand and analyze. The explanation uses concise, precise language; avoids redundant information; and structures content logically. Doctors can quickly grasp the core logic without additional effort.
- 5-7 points: Moderately easy to understand and analyze. The explanation is mostly clear but may contain minor redundancies or slightly complex sentence structures. Doctors can grasp the core logic with minimal effort, without needing to re-read repeatedly.
- 3-5 points: Difficult to understand and analyze. The explanation has confusing structure, ambiguous terminology, or excessive jargon. Doctors need to spend significant effort to understand the main content.
- 1-3 points: Nearly impossible to understand and analyze. The explanation is disorganized, uses inaccurate or obscure language, and contains massive redundant or irrelevant information. Doctors cannot effectively grasp the logic even after repeated reading.

A.7.2 LOGICAL ARGUMENTATION(LA)

The essence of an explanation lies in analyzing the reasoning process through logic, so the definition of Logical Argumentation is "Whether the expression is consistent and coherent, and whether the logic is clear and smooth".

- 7-10 points: Fully consistent, coherent, and logically rigorous. The explanation has a clear logical thread; each statement connects naturally to the next; there are no contradictions or logical gaps; and the reasoning process from premises to conclusions is complete and persuasive.
- 5-7 points: Mostly consistent, coherent, and logically clear. The overall logical thread is understandable, but there may be minor inconsistencies or weak transitions between statements. The reasoning process is generally complete with no major logical flaws.
- **3-5 points:** Inconsistent, incoherent, or logically confusing. The explanation has obvious logical gaps or occasional contradictions. The connection between statements is weak, making the overall logic difficult to follow.
- 1-3 points: Severely inconsistent, incoherent, or logically invalid. The explanation has serious contradictions; the reasoning process is chaotic or nonexistent; and there is no clear logical connection between statements, leading to complete loss of persuasiveness.

A.7.3 EVIDENCE-BASED MEDICINE(EBM)

The credibility of an explanation relies on the support of professional medical knowledge and evidence-based principles, so the definition of Evidence-Based Medicine is "Whether the explanation conforms to professional medical knowledge and evidence-based principles".

- 7-10 points: Fully conforms to professional medical knowledge and evidence-based principles.
 All medical claims are accurate and supported by well-recognized evidence. No medical errors or misinformation exist.
- 5-7 points: Mostly conforms to professional medical knowledge and evidence-based principles. Core medical claims are accurate, but minor details may lack strong evidence support or have slight imprecision. No critical medical errors.
- 3-5 points: Partially conforms to professional medical knowledge and evidence-based principles. There are noticeable medical inaccuracies or over-reliance on low-quality evidence. These issues do not completely invalidate the explanation but reduce its professional credibility.
- 1-3 points: Does not conform to professional medical knowledge and evidence-based principles. The explanation contains serious medical errors or promotes unsubstantiated claims. These issues make the explanation professionally unreliable.

A.7.4 COGNITIVE BIASING(CB)

When generating explanations, if only supporting factors are listed, it is prone to falling into the trap of intuitive judgment (Kahneman, 2011). Therefore, the definition of CB is "Whether the evidence listed in the explanation is comprehensive, encompassing both factors supporting and opposing the final prediction".

- 7-10 points: Extremely comprehensive evidence with no obvious bias. The explanation systematically lists key supporting factors and relevant opposing factors along with analysis. It also briefly discusses why opposing factors do not change the conclusion, showing balanced consideration.
- 5-7 points: Mostly comprehensive evidence with minimal bias. The explanation lists and simply analyze some key supporting factors and some major opposing factors. While a few minor opposing factors may be omitted, the overall presentation is balanced, and the bias is not obvious.
- **3-5 points:** Incomplete evidence with noticeable bias. The explanation focuses primarily on supporting factors and only mentions opposing factors superficially or omits important ones. The one-sided presentation makes the explanation lean heavily toward justifying the conclusion.
- 1-3 points: Highly incomplete evidence with severe bias. The explanation only lists factors supporting the final prediction and completely ignores all relevant opposing factors. It appears as a one-sided justification rather than a balanced explanation of the reasoning process.

A.8 AN EXPLANATION CASE

The patient has several risk factors contributing to catheterrelated thrombosis. The GLR of 1.82 is within a borderline range, with some studies suggesting values above 1.8 may reflect subclinical inflammation associated with increased thrombotic risk. Although D-dimer is normal at 0.1 mg/L, it may not reliably exclude CRT, especially in cancer patients where baseline levels can vary. The patient has lu known hypercoagulable condition that significantly increases VTE and CRT risk. He is undergoing thoracic the can cause vascular damage and immobility, further elevating thrombosis risk. The use of a PICC line adds mechanical endothelial injury as a contributory factor. Advanced age (74 years) is an independent risk factor for VTE due to decreased venous compliance and increased comorbidity burden. While platelet and hemoglobin levels are normal, and BMI is in the normal range, the combination of malignancy, PICC use, age, and thoracic therapy collectively heighten the likelihood of CRT despite the absence of che rapy and prior zation history.

> Cognitive Load: 7 Logical Argumentation: 6 Evidence-based Medicine: 7 Cognitive Biasing: 5

The patient's GLR of 0.73 is low, which may reflect a less pronounced inflammatory state and could be associated with a lower thrombotic risk, although the exact threshold for GLR in CRT remains debated. D-dimer level is 0.33 mg/L, well within the normal range, indicating low current fibrin turnover and minimal evidence of active thrombosis. The patient is receiving chemotherapy and has a PICC line, both of which are established risk factors for CRT due to endothelial injury and hypercoagulability. However, the patient is 46 years old, which confers a moderate age-related risk, and her platelet count (209.0 × 109/L) and hemoglobin (136.0 g/L) are within normal ranges, suggesting no significant thrombocytosis or anemia. BMI is 25.88, indicating overweight status, which is a mild risk factor for VTE. The patient is female with no prior catheterization or history of CRT, and has no known high-risk malignancies such as lung, gastric, or gynecologic cancers. Overall, despite the presence of chemotherapy and PICC, the absence of elevated D-dimer low GLR, and lack of other strong risk factors suggest a low likelihood of CRT

> Cognitive Load: 9 Logical Argumentation: 10 Evidence-based Medicine: 10 Cognitive Biasing: 9

Figure 4: Two samples of explanation generated are provided to help better understand the criteria used in explanation experiment. Protective factors in the text are highlighted in blue, while risk factors are highlighted in red.

Figure 4 presents the explanations generated for two patients respectively. The explanation on the left scored lower, with scores in the four dimensions (Cognitive Load, Logical Argumentation, Evidence-Based Medicine, Cognitive Biasing) being 7, 6, 7, and 5 in sequence. Specifically:

- *Cognitive Biasing*: Although the text mentions supportive evidence for thrombotic risk and supplements features associated with lower risk, basically meeting the requirement of evidence comprehensiveness, it insufficiently analyzes the opposing factors. It merely lists evidence of lower risk without providing valid information such as the association between this evidence and reduced CRT risk.
- Logical Argumentation: It generally follows a "total-subtotal-total" logic: first clarifying that the patient has multiple CRT risk factors, then analyzing each factor one by one, and finally concluding that "the superposition of multiple factors increases risk", with clear expression. However, there are deficiencies in the logical connection between some factors, and the argumentative relationship between certain features and "CRT risk" is not smooth enough.
- *Evidence-Based Medicine*: The overall prediction is based on evidence and conforms to professional medical knowledge, but when listing features associated with lower risk, it lacks professional support.

Based on the above three dimensions, although the text is generally logically coherent, with clear evidence-based core and evidence covering both positive and negative aspects, making it easy for doctors to read and understand; in the key part of the final summary, a large number of low-risk evidence contrary to the conclusion of "increased risk" are listed, which significantly increases the understanding pressure on doctors. Therefore, the overall Cognitive Load score is 7.

In contrast, the explanation on the right scored higher, with scores in the four dimensions being 9, 10, 10, and 9 in sequence:

- *Cognitive Biasing*: The text details two categories of features—"those increasing CRT risk" and "those reducing CRT risk", conducts in-depth analysis of each feature, and does not omit key influencing factors, with comprehensive and balanced evidence.
- *Logical Argumentation*: It adopts a "subtotal-total" logic for argumentation, clearly explaining how each feature directly or indirectly affects CRT risk, and finally draws the conclusion that "CRT risk is reduced", with a clear and coherent argumentation process.
- Evidence-Based Medicine: The analysis of risk factors conforms to recognized evidence-based conclusions, and the analysis of protective factors (i.e., factors reducing risk) also meets clinical testing standards; at the same time, it specifically mentions the controversy that "the threshold of GLR for CRT diagnosis has not been clearly defined", which not only respects evidence-based principles but also does not affect the professionalism of the overall argumentation.

Overall, the understanding cost for professional doctors to read this text is extremely low. They can quickly grasp the core argument framework without the need to additionally verify the professionalism of the information, and there is basically no cognitive burden. Therefore, the Cognitive Load score is 9.

A.9 MAIN EXPERIMENT RESULTS

This section supplements Section 4.2. Table 2 presents the predictive performance of all methods in the main experiment. As can be seen from the table, RCA+Qwen2.5 and RCA+GPT-4.1 have nearly outperformed all baselines across the three metrics on the three datasets. Notably, on the heart disease dataset, the accuracy of RCA+GPT-4.1 is 20% higher than that of the top-performing baseline, accompanied by excellent MCC and F1-score. It is worth noting that LLM-based methods generally perform poorly on the CRT dataset and the Heart Disease dataset. However, o3-mini+Code achieves promising results on the CRT dataset, which indicates that code assistance enhances the reasoning ability of the o3-mini.

Table 3-5 present the explanation scores of all methods across the three datasets. As can be observed from the tables, the RCA methods outperform all baselines and achieve the highest scores across every dataset. This result indicates that the rules derived from in-depth understanding and summarization of data can effectively enhance the quality of interpretation. Among the four metrics, "Cognitive Biasing" consistently yields the lowest scores. This phenomenon suggests that when generating explanations, LLMs tend to prioritize listing factors that support their own conclusions—a characteristic that also aligns with the intuitive behavioral patterns of humans.

Table 2: Accuracy, MCC and F1-score results in main experiment. RCA achieve almost best performance across all datasets, with Accuracy and MCC scores that rival all those of tree-based methods known for their effectiveness with tabular data. Moreover, it significantly outperforms LLM and LLM-based agents.

	Datasets	CRT			Diabetes			Н	Heart Disease		
	Algorithm	Acc	MCC	F1-score	Acc	MCC	F1-score	Acc	MCC	F1-score	
Traditional	Lasso	0.5238	0.1261	0.1667	0.6867	0.3596	0.2857	0.5337	-0.0065	0.1212	
MLs	Catboost	0.7460	0.1472	0.2000	0.7590	0.5252	0.3478	0.5440	0.0182	0.1212	
	Qwen2.5-72B	0.2698	-0.0073	0.1154	0.7470	0.4244	0.3333	0.3101	-0.0870	0.0851	
LLM-based	DeepSeekV3	0.2698	0.0111	0.1154	0.5542	0.4797	0.2222	0.2228	-0.0383	0.1091	
Method	DeepSeekV3.1	0.1587	0.1181	0.1212	0.7229	0.3687	0.5660	0.2073	-0.1877	0.3377	
	GPT-4.1	0.1746	-0.0797	0.1034	0.7470	0.4187	0.3333	0.2280	-0.0502	0.1091	
	DeepSeek-R1	0.1429	0.0805	0.1290	0.7229	0.3615	0.3200	0.2124	-0.1367	0.0741	
	o3-mini	0.3175	0.0232	0.1224	0.5060	-0.1352	0.0606	0.2176	-0.0996	0.1091	
Reasoning	o4-mini	0.2381	0.1220	0.1455	0.7470	0.4173	0.4878	0.2280	-0.0114	0.3550	
LLMs	Qwen3-30B	0.1270	0.0678	0.1270	0.7470	0.4187	0.5882	0.2228	-0.0353	0.3534	
	Qwen3-235B	0.2063	0.1063	0.1380	0.7108	0.3451	0.5556	0.2280	-0.1430	0.3258	
	GPT-5	0.2857	0.0024	0.1176	0.7470	0.4106	0.5333	0.2280	-0.1602	0.3196	
	Qwen2.5-72B+Tools	0.2857	0.0023	0.1176	0.7470	0.4244	0.3333	0.3161	-0.0797	0.0851	
	DeepSeekV3+Tools	0.2381	0.1242	0.1429	0.7470	0.4187	0.3333	0.2280	0.0543	0.1404	
	DeepSeekV3.1+Tools	0.1905	-0.0650	0.1053	0.7349	0.4002	0.5926	0.2228	-0.0352	0.3534	
	GPT-4.1+Tools	0.2698	-0.0073	0.1154	0.7590	0.4454	0.3478	0.2694	-0.1203	0.0800	
LLM-based	DeepSeek-R1+Tools	0.1111	0.0582	0.1250	0.7349	0.3873	0.3200	0.2435	0.0351	0.1111	
Agent	o3-mini + Tools	0.2063	-0.0577	0.1071	0.4359	-0.2441	0.0000	0.2265	-0.1003	0.1091	
	o4-mini+Tools	0.2381	0.1220	0.1455	0.7108	0.3169	0.4783	0.2228	-0.0721	0.3478	
	Qwen3-30B+Tools	0.0952	0.0471	0.1231	0.7108	0.3169	0.4783	0.2073	-0.1877	0.3377	
	Qwen3-235B+Tools	0.1905	0.0946	0.1091	0.6987	0.3223	0.5455	0.2073	-0.2282	0.3139	
	o3-mini+Code	0.5079	0.1179	0.1622	0.6747	0.3861	0.2857	0.2850	-0.0594	0.1176	
RCA	RCA+Qwen2.5	0.7778	0.1739	0.2222	0.7831	0.5406	0.7097	0.5647	0.0547	0.1290	
	RCA+GPT-4.1	0.8730	0.1373	0.2000	0.7470	0.4244	0.6038	0.7461	0.1493	0.2898	

A.10 ROBUST EXPERIMENT RESULTS

Table 6 presents the predictive performance across all datasets, while Tables 7-9 respectively show the explanation scores for each of the three datasets. As can be seen from the table data, RCA+GPT-4.1 achieves the best performance in nearly all predictive metrics while maintaining competitive explanation scores, with RCA+Qwen2.5 performing slightly worse in prediction task. Such minor fluctuations indicate that RCAenables LLMs to maintain robustness against data noise. Furthermore, a comprehensive analysis of the aforementioned tables reveals no direct correlation between predictive performance and explanation scores: for instance, Qwen3-235B, which performs poorly in terms of accuracy, achieves impressive scores across all four explanation metrics. This phenomenon demonstrates that modern LLMs possess strong capabilities in generating explanatory texts, regardless of whether their predictive results are correct or not.

A.11 ABLATION STUDY

As shown in Table 10, removing any module from RCA leads to a significant drop in the four explanation scores, which indicates that each module plays an irreplaceable role in the process of explanatory reasoning. Among these, the removal of the rules check module causes the largest drop in scores. This phenomenon may be attributed to the fact that without the inspection module, the quality of rules becomes difficult to control effectively, which in turn negatively impacts the quality of explanatory texts.

Table 3: Explanation experiment results on CRT dataset

CRT		CL	LA	EBM	CB
Traditional	Lasso	7.60	8.33	7.31	7.21
MLs	Catboost	7.39	8.13	7.31	7.08
	Qwen2.5-72B	7.92	8.16	8.25	7.14
LLM-Based	DeepSeek-V3	7.98	8.27	8.22	7.19
Methods	DeepSeek-V3.1	7.41	8.11	8.13	6.73
	GPT-4.1	7.79	8.25	8.30	7.31
	DeepSeek-R1	7.82	7.97	8.27	7.25
	o3-mini	6.84	7.22	6.24	6.19
Reasoning	o4-mini	7.56	8.06	8.25	6.81
LLMs	Qwen3-30B	7.87	8.49	8.41	7.06
	Qwen3-235B	8.22	8.81	8.71	7.49
	GPT-5	7.63	8.22	8.40	7.05
	Qwen2.5-72B+Tools	7.87	8.35	8.29	7.16
	DeepSeek-V3.1+Tools	7.87	8.41	8.35	6.51
	GPT-4.1+Tools	7.89	8.30	8.32	7.35
LLM-Based	DeepSeek-R1+Tools	7.87	8.00	8.17	7.14
Agents	o3-mini+Tools	6.22	7.30	6.20	6.16
Agents	o4-mini+Tools	7.29	7.87	8.05	5.73
	Qwen3-30B+Tools	7.79	8.63	8.41	6.57
	Qwen3-235B+Tools	8.13	8.86	8.86	7.11
	o3-mini+Code	6.93	7.60	6.68	6.17
RCA	RCA + Qwen2.5	8.24	8.89	8.47	7.61
NCA	RCA + GPT-4.1	8.16	8.59	8.87	7.62

Table 4: Explanation experiment results on Diabetes dataset

Diabetes		CL	LA	EBM	CB
Traditional	Lasso	7.83	8.39	8.33	6.24
MLs	Catboost	7.71	8.27	8.19	6.12
	Qwen2.5-72B	7.61	8.20	8.10	5.98
LLM-Based	DeepSeek-V3	7.55	8.04	8.29	5.99
Methods	DeepSeek-V3.1	7.85	8.23	8.12	6.04
	GPT-4.1	7.84	8.14	8.33	5.85
	DeepSeek-R1	7.73	8.12	7.99	5.88
	o3-mini	7.67	8.12	8.16	5.89
Reasoning	o4-mini	7.82	8.13	8.08	6.13
LLMs	Qwen3-30B	7.94	8.43	8.51	6.72
	Qwen3-235B	7.84	8.41	8.31	6.30
	GPT-5	7.66	8.24	8.53	6.24
	Qwen2.5-72B+Tools	7.86	8.29	8.66	6.43
	DeepSeek-V3+Tools	7.57	8.09	8.22	6.03
	DeepSeek-V3.1+Tools	7.83	8.23	8.55	6.41
	GPT-4.1+Tools	7.81	8.25	8.01	6.12
LLM-Based	DeepSeek-R1+Tools	7.73	8.24	8.11	5.93
Agents	o3-mini+Tools	7.62	8.25	8.13	5.90
	o4-mini+Tools	7.85	8.27	8.12	6.03
	Qwen3-30B+Tools	7.82	8.13	8.08	6.13
	Qwen3-235B+Tools	7.88	8.33	8.11	6.38
	o3-mini+Code	7.17	7.60	6.76	5.49
RCA	RCA + Qwen2.5	8.13	8.57	8.74	6.43
KCA	RCA + GPT-4.1	8.03	8.38	8.63	6.94

A.12 PROMPT TEMPLATES

In this section we will provide all prompt templates used in RCA.

Table 5: Explanation experiment results on Heart Disease dataset

Heart Disease		CL	LA	EBM	CB
Traditional	Lasso	7.60	8.27	8.63	5.94
MLs	Catboost	7.50	8.21	8.72	5.94
	Qwen2.5-72B	7.63	8.33	8.77	6.14
LLM-Based	DeepSeek-V3	7.52	8.25	8.62	6.05
Methods	DeepSeek-V3.1	7.35	8.04	8.55	5.69
	GPT-4.1	7.45	8.08	8.64	5.77
	DeepSeek-R1	7.40	7.98	8.59	5.75
	o3-mini	7.33	7.89	8.84	5.46
Reasoning	o4-mini	7.31	7.83	8.73	5.38
LLMs	Qwen3-30B	7.41	8.06	8.49	5.82
	Qwen3-235B	7.67	8.34	8.84	6.36
	GPT-5	7.36	8.07	8.59	5.62
	Qwen2.5-72B+Tools	7.63	8.34	8.74	6.0
	DeepSeek-V3+Tools	7.43	8.21	8.55	6.11
	DeepSeek-V3.1+Tools	7.37	7.97	8.46	5.72
	GPT-4.1+Tools	7.42	8.01	8.58	5.70
LLM-Based	DeepSeek-R1+Tools	7.38	7.87	8.55	5.73
Agents	o3-mini+Tools	7.31	7.90	8.76	5.48
	o4-mini+Tools	7.35	7.96	8.51	5.66
	Qwen3-30B+Tools	7.39	7.94	8.56	5.73
	Qwen3-235B+Tools	7.58	8.28	8.91	6.20
	o3-mini+Code	7.28	7.88	8.67	5.38
RCA	RCA + Qwen2.5	7.62	8.47	8.94	6.18
NCA	RCA + GPT-4.1	7.74	8.53	8.79	6.42

A.12.1 RULES OPTIMIZATION

Table 11 13 shows the prompt template for M_{ref} to iteratively extract rules in the self-reflection process. In the prompt, incorrectly predicted samples, along with previous rules and data distribution are fed to M_{ref} , including features and true labels. Then M_{ref} will consider what caused the wrong predictions and optimized the rule base. If the previous rule is empty, it means extracting initial rules. We specifically emphasized adherence to medical knowledge in the prompts and incorporated negative example to standardize rule generation.

A.12.2 RULES CHECK

Table 14-16 shows the prompt template for M_{chk} to check and delete the rules. At the end of each epoch, M_{chk} checks the rule base to maintain the quality of the rules. The prompt lists several major errors we have identified and provides example rules tailored to specific diseases and features. We can see that the prompt also include previous distribution and previous rules. It instructs M_{chk} to examine each rule and remove incorrect or low-quality rule, preventing them from affecting predictions.

A.12.3 DISEASE PREDICTION

Table 17-19 show the prompt template for M_{pred} to generate prediction and explanation for the patient. We can see that the prompts are largely consistent across the three datasets, with only minor differences in wording. The prompt on all three datasets contains both positive and negative examples to provide demonstrations for the M_{pred} .

Table 6: Accuracy, MCC and F1-score results in robust experiment. RCA achieve almost best performance across all datasets, with Accuracy and MCC scores that rival all those of tree-based methods known for their effectiveness with tabular data. Moreover, it significantly outperforms LLM and LLM-based agents.

	Datasets		w/o GLR			Missin	g		Abnormal		
	Algorithm	Acc	MCC	F1-score	Acc	MCC	F1-score	Acc	MCC	F1-score	
Traditional	Lasso	0.5238	0.1261	0.1667	0.5079	0.1179	0.1622	0.6667	0.0927	0.1600	
MLs	Catboost	0.6825	0.1021	0.1667	0.5079	0.0041	0.1143	0.5238	0.0124	0.1176	
	Qwen2.5	0.1277	0.0678	0.1270	0.3333	-0.0921	0.0870	0.3175	-0.1021	0.0851	
LLM-based	DeepSeekV3	0.2539	-0.0152	0.1132	0.2857	0.0048	0.1176	0.2063	-0.1082	0.1071	
Method	DeepSeekV3.1	0.2063	-0.0517	0.1071	0.1429	-0.1151	0.1000	0.2073	-0.1877	0.3377	
	GPT-4.1	0.2381	-0.0280	0.1111	0.2222	-0.0395	0.1091	0.1905	0.0993	0.1356	
	DeepSeek-R1	0.0793	0.0331	0.1212	0.1429	0.0764	0.1290	0.0635	0.0000	0.0000	
	o3-mini	0.3968	-0.0315	0.0952	0.3175	0.1001	0.1569	0.2063	0.0971	0.1379	
Reasoning	o4-mini	0.2222	-0.0394	0.1091	0.2222	0.1151	0.1429	0.2280	-0.0114	0.3550	
LLMs	Qwen3-30B	0.1270	0.0678	0.1270	0.1269	0.0678	0.1270	0.2228	-0.0353	0.3534	
	Qwen3-235B	0.1587	0.0844	0.1311	0.2063	0.1063	0.1379	0.2280	-0.1430	0.3258	
	GPT-5	0.3333	-0.0921	0.0870	0.3333	0.0293	0.1250	0.2381	-0.0281	0.1111	
	Qwen2.5-72B+Tools	0.2539	0.1263	0.1455	0.3175	-0.1021	0.0851	0.2222	0.1131	0.1404	
	DeepSeekV3+Tools	0.1746	-0.2155	0.0714	0.2222	0.1267	0.1404	0.1905	0.1089	0.1356	
	DeepSeekV3.1+Tools	0.1905	0.0650	0.1053	0.2540	0.1263	0.1455	0.1429	0.0764	0.1290	
	GPT-4.1+Tools	0.2222	-0.0395	0.1091	0.2857	-0.1235	0.0816	0.1905	-0.0650	0.1053	
LLM-based	DeepSeek-R1+Tools	0.1111	0.0582	0.1250	0.1746	0.0921	0.1333	0.0793	0.0331	0.1212	
Agent	o3-mini + Tools	0.3175	-0.1099	0.0851	0.2063	0.0971	0.1379	0.1746	-0.6240	0.0851	
	o4-mini+Tools	0.2381	-0.0281	0.1111	0.2222	0.1131	0.1404	0.2063	0.0989	0.1091	
	Qwen3-30B+Tools	0.1587	0.0845	0.1311	0.1429	0.0764	0.1290	0.1746	0.0921	0.1333	
	Qwen3-235B+Tools	0.1429	0.0764	0.1290	0.2381	0.1198	0.1429	0.1905	0.0993	0.1356	
	o3-mini+Code	0.4603	-0.1343	0.0556	0.5873	0.0462	0.1333	0.4286	-0.0919	0.1000	
RCA	RCA+Qwen2.5	0.7143	0.1235	0.1818	0.6984	0.1126	0.1739	0.7302	0.1350	0.1904	
KCA	RCA+GPT-4.1	0.6507	0.1979	0.2143	0.8889	0.1644	0.2222	0.8254	0.2232	0.2666	

Table 11: Prompt template for reflective rules extraction on CRT dataset

You are an advanced reasoning agent that can improve based on self reflection. The original task is "Given clinical features of tumor patient, estimate whether the patient has the catheter related thrombosis(CRT) or not and explain your reasoning.". Now you will be given the previous rules and some wrong samples that you have attempted to predict CRT but failed. Considering patients' clinical features and their true CRT results, you need to reflect on and revise rules to help CRT prediction.

The rules must be supported by medical knowledge. Note that rules like "If the patient has a BMI value between 10 and 50 and a history of previous CRT, predict no catheter-related thrombosis." are not reasonable — you cannot predict CRT only based on BMI and previous CRT history.

There might be some outliers in the data. You will also be given the features distribution on the whole dataset. You can determine the outliers based on distribution before summarizing the rules. Don't utilize relationship between outliers and CRT, and don't let your rules be destroyed by outliers easily. Don't exclude patients with outliers, you should use other features to predict CRT for them.

Keep the rules brief. Only output rules. Your rules must be general enough for any patients. Give your response in this format:

Rules, which should be a list of rules, each rule is a short sentence.

Previous distribution: {distribution}

Previous rules: {rules}

```
(If it is empty, it means summarizing the initial rules)
Wrong samples:
{samples}
Rules:
```

Table 12: Prompt template for reflective rules extraction on Diabetes dataset

You are an advanced reasoning agent that can improve based on self reflection. The original task is "Given clinical features of patient, estimate whether the patient has the diabetes or not and explain your reasoning.". Now you will be given the previous rules and some wrong samples that you have attempted to predict diabetes but failed. Considering patients' clinical features and their true diabetes results, you need to reflect on and revise rules to help diabetes prediction.

The rules must be supported by medical knowledge. Note that rules like "If the patient has a BMI value greater than 25, predict diabetes." are not reasonable - you cannot predict diabetes only based on BMI.

There might be some outliers in the data. You will also be given the features distribution on the whole dataset. You can determine the outliers based on distribution before summarizing the rules. Don't utilize relationship between outliers and diabetes, and don't let your rules be destroyed by outliers easily. Don't exclude patients with outliers, you should use other features to predict diabetes for them.

Keep the rules brief. Only output rules. Your rules must be general enough for any patients. Give your response in this format:

Rules, which should be a list of rules, each rule is a short sentence.

```
Data distribution:
{distribution}

Previous rules:
{rules}
(If it is empty, it means summarizing the initial rules)

Wrong samples:
{samples}
Rules:
```

Table 13: Prompt template for reflective rules extraction on Heart Disease dataset

You are an advanced reasoning agent that can improve based on self reflection. The original task is "Given clinical features of patient, estimate whether the patient has the heart disease or not and explain your reasoning.". Now you will be given the previous rules and some wrong samples that you have attempted to predict heart disease but failed. Considering patients' clinical features and their true heart disease results, you need to reflect on and revise rules to help heart disease prediction.

The rules must be supported by medical knowledge. Note that rules like "If the patient has a blood pressure greater than 160 and has diabetes, predict no heart disease." are not reasonable — you cannot predict heart disease only based on blood pressure and diabetes.

There might be some outliers in the data. You will also be given the features distribution on the whole dataset. You can determine the outliers based on distribution before summarizing the rules. Don't utilize relationship between outliers and heart disease, and don't let your rules be destroyed by outliers easily. Don't exclude patients with outliers, you should use other features to predict heart disease for them.

Keep the rules brief. Only output rules. Your rules must be general enough for any patients. Give your response in this format:

Rules, which should be a list of rules, each rule is a short sentence.

```
Data distribution:
{distribution}

Previous rules:
{rules}
(If it is empty, it means summarizing the initial rules)

Wrong samples:
{samples}
Rules:
```

Table 14: Prompt template for rules check on CRT dataset

You are an advanced reasoning agent that can improve based on self reflection. The original task is "Given clinical features of tumor patient and some prediction rules, estimate whether the patient has the catheter related thrombosis(CRT) or not and explain your reasoning" Given the previous rules and the features distribution, you need to check and delete the error rules.

Rules like "If the patient has a BMI value between 10 and 50 and a history of previous CRT, predict no catheter-related thrombosis." are not reasonable, because it's inconsistent with medical knowledge – you cannot predict CRT only based on BMI and previous CRT history.

Also, there might be some outliers in data. Rules that utilize relationship between outliers and disease, like "If the patient has a D-dimer level between 0.1 and 0.79, but any numerical feature is an extreme outlier, they are less likely to develop CRT." is forbidden. However, outliers could mislead prediction, so you should indicate in the rules how to identify outliers. Don't exclude patients with outliers, you should use other features to support disease prediction for them.

And rules that are too specific for certain patient are awful. You need to delete rules similar to those listed above. Give your response in this format:

Rules, which should be a list of rules, each rule is a short sentence.

```
Previous distribution: {distribution}

Previous rules: {rules}
```

Rules:

Table 7: Explanation experiment results on CRT dataset w/o GLR

		CL	LA	EBM	CB
Traditional	Lasso	7.79	8.52	8.27	6.94
MLs	Catboost	7.78	8.48	8.41	6.82
	Qwen2.5-72B	7.68	8.41	8.32	6.43
LLM-Based	DeepSeek-V3	7.71	8.41	8.42	6.77
Methods	DeepSeek-V3.1	7.83	8.56	8.48	6.71
	GPT-4.1	7.71	8.43	8.52	6.51
	DeepSeek-R1	7.98	8.76	8.63	6.82
	o3-mini	7.56	8.22	8.27	6.13
Reasoning	o4-mini	7.37	8.08	8.21	6.03
LLMs	Qwen3-30B	7.91	8.49	8.67	6.76
	Qwen3-235B	8.22	8.75	8.86	7.12
	GPT-5	7.63	8.44	8.14	6.81
	Qwen2.5-72B+Tools	7.76	8.48	8.56	6.52
	DeepSeek-V3+Tools	7.73	8.52	8.37	6.73
	DeepSeek-V3.1+Tools	7.84	8.49	8.51	6.69
	GPT-4.1+Tools	7.73	8.56	8.76	6.42
LLM-Based	DeepSeek-R1+Tools	8.11	8.79	8.70	6.84
Agents	o3-mini+Tools	7.46	8.05	7.98	6.05
	o4-mini+Tools	7.49	8.21	8.22	6.41
	Qwen3-30B+Tools	7.78	8.56	8.57	6.49
	Qwen3-235B+Tools	8.27	8.92	9.06	7.16
	o3-mini+Code	7.90	8.54	8.68	6.67
RCA	RCA + Qwen2.5	8.35	8.97	8.99	7.26
11011	RCA + GPT-4.1	8.28	8.75	9.13	7.21

Table 15: Prompt template for rules check on Diabetes dataset

You are an advanced reasoning agent that can improve based on self reflection. The original task is "Given clinical features of patient and some prediction rules, estimate whether the patient has the diabetes or not and explain your reasoning" Given the previous rules and the features distribution, you need to check and delete the error rules.

Rules like "If the patient has a BMI value greater than 25, predict diabetes." are not reasonable, because it's inconsistent with medical knowledge — you cannot predict diabetes only based on BMI.

Also, there might be some outliers in data. Rules that utilize relationship between outliers and disease, like "If the patient has a Diastolic blood pressure between 80 mmHg and 90 mmHg, but any numerical feature is an extreme outlier, they are less likely to develop diabetes." is forbidden. However, outliers could mislead prediction, so you should indicate in the rules how to identify outliers. Don't exclude patients with outliers, you should use other features to support disease prediction for them.

And rules that are too specific for certain patient are awful. You need to delete rules similar to those listed above. Give your response in this format:

Rules, which should be a list of rules, each rule is a short sentence.

Previous distribution:
{distribution}

Previous rules:
{rules}

Rules:

Table 16: Prompt template for rules check on Heart Disease dataset

You are an advanced reasoning agent that can improve based on self reflection. The original task is "Given clinical features of tumor patient and some prediction rules, estimate whether the patient has the heart disease or not and explain your reasoning" Given the previous rules and the features distribution, you need to check and delete the error rules.

Rules like "If the patient has a blood pressure greater than 160 and has diabetes, predict no heart disease." are not reasonable, because it's inconsistent with medical knowledge — you cannot predict heart disease only based on blood pressure and diabetes.

Also, there might be some outliers in data. Rules that utilize relationship between outliers and disease, like "If the patient has a CRP level between 10 and 12, but any numerical feature is an extreme outlier, they are less likely to develop heart disease." is forbidden. However, outliers could mislead prediction, so you should indicate in the rules how to identify outliers. Don't exclude patients with outliers, you should use other features to support disease prediction for them.

And rules that are too specific for certain patient are awful. You need to delete rules similar to those listed above. Give your response in this format:

Rules, which should be a list of rules, each rule is a short sentence.

Previous distribution:
{distribution}

Previous rules:
{rules}

Rules:

Table 17: Prompt template for disease prediction on CRT dataset

Given clinical features of tumor patient, estimate whether the patient has the catheter related thrombosis(CRT) or not and explain your reasoning. You will be given some rules for prediction and distribution of training dataset. You can refer to the following rules, but don't limit yourself to them. Remember there are some outliers in the data. Give your response in this format:

- (1) CRT Prediction, which should be either "no catheter-related thrombosis" or "catheter-related thrombosis".
- (2) Explanation, which should be in a single, short paragraph.

Here are some examples:

Features:Granulocyte-to-lymphocyte ratio is 1.44, D-dimer is 0.19, chemotherapy, catheterization is CVC(Central Venous Catheter), no thoracic therapy, age at hospital is 29, platelet is 353.0, hemoglobin is 138.0, BMI is 18.83, gender is male, history of previous catheterization, no history of previous catheter related thrombosis, no lung cancer, no gastric cancer, lymphoma, no gynecologic tumors, no urologic tumors.

CRT Prediction: no catheter-related thrombosis

Explanation: The GLR of the patients was 1.44. Some studies have suggested that GLR is associated with thrombosis, but the normal threshold value of GLR is still under debate and the range of normal values of GLR varies from study to study. The patient's GLR value of 1.44 was within the normal range, which may indicate a lower risk of thrombosis. The patient's D-dimer level was 0.19 mg/L, which is within the normal range (less than 0.5 mg/L is generally considered normal), and lower D-dimer levels are usually associated with lower thrombotic risk. The patient is receiving chemotherapy, a known risk factor for VTE. Chemotherapy patients have a 6.5-fold elevated risk of thrombosis. The patient is using a CVC (Central Venous Catheter). The use of a central venous catheter is itself a risk factor for VTE, especially in oncology patients. The patient was relatively young at 29 years of age, and usually younger patients have a lower risk of thrombosis. The patient's platelet level was 353.0, slightly above the normal range, and thrombocytosis is a predictor of VTE. The patient's hemoglobin level was 138.0 g/L, which is in the normal range.BMI: The patient's BMI was 18.83, which is in the underweight range, and it is generally accepted that higher BMIs are more likely to result in CRT. The patient was male, and the effect of gender on thrombotic risk has varied in different studies. The patient had a history of previous catheterization, which may increase the risk of CRT. The patient had no history of prior thrombosis and did not develop certain tumors, which may indicate a lower risk of thrombosis. In summary, the patient was predicted to have a low risk of CRT.

Features:Granulocyte-to-lymphocyte ratio is 2.73, D-dimer is 0.1, chemotherapy, catheterization is PICC(Peripherally Inserted Central Catheter), no thoracic therapy, age at hospital is 30, platelet is 267.0, hemoglobin is 108.0, BMI is 26.04, gender is female, no history of previous catheterization, no history of previous catheter related thrombosis, no lung cancer, no gastric cancer, no lymphoma, no gynecologic tumors, no urologic tumors.

CRT Prediction: catheter-related thrombosis

Explanation:GLR is an indicator of inflammation and immune status.GLR 2.73 is a relatively high value and may indicate the presence of an inflammatory response, which may be associated with an increased risk of thrombosis.D-dimer is a marker of coagulation and fibrinolysis.A level of 0.1 is usually considered normal or only slightly elevated and is not sufficient to directly diagnose VTE. Therefore, this level of D-dimer is unlikely to indicate the presence of CRT. chemotherapy may increase a patient's coagulation status because it can cause vascular endothelial injury and inflammation, which can increase the risk of thrombosis. the use of a PICC is a known risk factor for CRT because catheters can cause vascular endothelial injury and inflammation, which can promote thrombosis. Younger age is associated with a relatively lower risk of VTE. Platelet counts above the normal range may indicate a risk of inflammation or thrombosis. A slightly lower hemoglobin level may indicate mild anemia, but this level usually does not directly affect the risk of thrombosis. A slightly higher body mass index (BMI) indicates that the patient may be overweight, which is a risk factor for VTE. Gender is not an independent risk factor for CRT. There was no history of previous catheterization, which reduced the patient's risk of CRT. There is no history of catheter-related thrombosis, which reduces the patient's risk of CRT. No history of certain malignancies, which are known risk factors for VTE and CRT. In summary, the patient's risk of having catheter-related thrombosis is relatively high.

```
(END OF EXAMPLES)
Here are some rules:
{rules}
(If it is empty, it means there is no rule.)
(END OF RULES)
```

```
Here is the distribution:
{distribution}
(END OF DISTRIBUTION)

Features:
{features}

CRT Prediction:
```

Table 18: Prompt template for disease prediction on Diabetes dataset

Given clinical features of patient, estimate whether the patient has the diabetes or not and explain your reasoning. You will be given some rules for prediction and distribution of training dataset. You can refer to the following rules, but don't limit yourself to them. Remember there are some outliers in the data. Give your response in this format:

- (1) Diabetes Prediction, which should be either "no diabetes" or "diabetes".
- (2) Explanation, which should be in a single, short paragraph.

Here are some examples:

Features: Number of pregnancies is 1, Plasma glucose concentration (2-hour test) level is 135, Diastolic blood pressure is 54 mm Hg, Triceps skin fold thickness is 0 mm, 2-Hour serum insulin level is 0 mu U/ml, BMI is 26.7, DiabetesPedigreeFunction(Genetic diabetes score) is 0.687, Age is 62.

Diabetes Prediction: no diabetes

Explanation: The patient's 2-hour plasma glucose level is 135 mg/dL, which is below the diagnostic threshold for diabetes (>=200 mg/dL) and even below the range for prediabetes (140-199 mg/dL). While factors like age (62), overweight BMI (26.7), and a moderate genetic risk score (0.687) increase diabetes risk, the absence of elevated glucose levels within diagnostic ranges and other features (e.g., low triceps skin fold thickness, low insulin level) do not meet criteria for diabetes. Diagnosis primarily relies on glucose levels, which here are within normal limits.

Features: Number of pregnancies is 4, Plasma glucose concentration (2-hour test) level is 171, Diastolic blood pressure is 72 mm Hg, Triceps skin fold thickness is 0 mm, 2-Hour serum insulin level is 0 mu U/ml, BMI is 43.6, DiabetesPedigreeFunction(Genetic diabetes score) is 0.479, Age is 26.

Diabetes Prediction: diabetes

Explanation: The patient's plasma glucose concentration (171 mg/dL) exceeds the prediabetes threshold (>=140 mg/dL) and approaches the diabetes range, combined with a markedly elevated BMI (43.6, class III obesity), a major risk factor for type 2 diabetes. The genetic risk score (0.479) and history of 4 pregnancies (potential gestational diabetes risk) further support this prediction. While triceps skinfold thickness and insulin levels of 0 suggest possible data anomalies, the high glucose and BMI strongly indicate diabetes likelihood despite the patient's younger age (26).

```
(END OF EXAMPLES)
Here are some rules:
{rules}
(If it is empty, it means there is no rule.)
(END OF RULES)
```

Here is the distribution:
{distribution}
(END OF DISTRIBUTION)

Features: {features}

Diabetes Prediction:

Table 19: Prompt template for disease prediction on Heart Disease dataset

Given clinical features of tumor patient, estimate whether the patient has the heart disease or not and explain your reasoning. You will be given some rules for prediction and distribution of training dataset. You can refer to the following rules, but don't limit yourself to them. Remember there are some outliers in the data. Give your response in this format:

- (1) Heart Disease Prediction, which should be either "no heart disease" or "heart disease".
- (2) Explanation, which should be in a single, short paragraph.

Here are some examples:

Features: Age is 62.0, Gender is Female, Blood Pressure is 133.0, Cholesterol Level is 166.0, Exercise Habits is Medium, Smoking is No, Family Heart Disease is No, Diabetes is No, BMI is 25.739170533963147, High Blood Pressure is No, Low HDL Cholesterol is Yes, High LDL Cholesterol is No, Stress Level is Low, Sleep Hours is 5.493276805328829, Sugar Consumption is Medium, Triglyceride Level is 126.0, Fasting Blood Sugar is 102.0, CRP Level is 11.60991435489297, Homocysteine Level is 8.297757016065253

Heart disease Prediction: heart disease

Explanation: The patient has several risk factors for heart disease. At 62 years old, the patient has a cholesterol level of 166, and despite having normal blood pressure according to the "High Blood Pressure" marker, a blood pressure of 133 is relatively close to the elevated range. The presence of low HDL cholesterol is a risk factor for heart disease. The C-reactive protein (CRP) level of 11.60991435489297 is elevated, indicating possible inflammation in the body, which is associated with heart disease. Although the patient has a medium level of exercise and no family history of heart disease or diabetes, the combination of age, low HDL cholesterol, and elevated CRP level increases the likelihood of having heart disease.

Features: Age is 35.0, Gender is Male, Blood Pressure is 159.0, Cholesterol Level is 261.0, Exercise Habits is Low, Smoking is No, Family Heart Disease is No, Diabetes is Yes, BMI is 21.63849835899007, High Blood Pressure is No, Low HDL Cholesterol is Yes, High LDL Cholesterol is No, Stress Level is High, Sleep Hours is 4.296875738592791, Sugar Consumption is Medium, Triglyceride Level is 385.0, Fasting Blood Sugar is 136.0, CRP Level is 1.9462702594315329, Homocysteine Level is 11.140952179886469

Heart disease Prediction: no heart disease

Explanation: Although the patient presented with multiple risk factors such as elevated blood pressure, high cholesterol levels, diabetes, high triglycerides, high stress, low sleep hours, elevated CRP, and low HDL cholesterol, it has been determined that he has no heart disease. It is possible that there are mitigating factors not mentioned, such as effective medical management or significant lifestyle changes that reduce the impact of these risk factors on the heart.

```
(END OF EXAMPLES)

Here are some rules:
{rules}
(If it is empty, it means there is no rule.)
(END OF RULES)

Here is the distribution:
{distribution}
(END OF DISTRIBUTION)

Features:
{features}
Heart Disease Prediction:
```

Table 8: Explanation experiment results on CRT dataset missing 10%

		CL	LA	EBM	CB
Traditional	Lasso	7.89	8.24	8.13	6.90
MLs	Catboost	7.91	8.33	8.27	6.93
	Qwen2.5-72B	7.69	8.37	8.02	6.39
LLM-Based	DeepSeek-V3	7.65	8.31	8.11	6.28
Methods	DeepSeek-V3.1	7.78	8.29	8.17	6.32
	GPT-4.1	7.52	8.22	8.19	6.52
	DeepSeek-R1	7.79	8.24	8.21	6.75
	o3-mini	7.33	7.92	7.71	5.87
Reasoning LLMs	o4-mini	7.38	7.95	7.76	5.97
	Qwen3-30B	7.68	8.41	8.21	6.59
	Qwen3-235B	7.95	8.65	8.70	7.03
	GPT-5	7.82	8.30	8.14	6.73
	Qwen2.5-72B+Tools	7.63	8.25	8.05	6.38
	DeepSeek-V3+Tools	7.58	8.22	8.17	6.08
	DeepSeek-V3.1+Tools	7.72	8.37	8.13	6.33
	GPT-4.1+Tools	7.60	8.37	8.26	6.49
LLM-Based	DeepSeek-R1+Tools	7.62	8.35	8.25	6.67
Agents	o3-mini+Tools	7.41	8.07	7.81	6.03
	o4-mini+Tools	7.44	7.94	8.14	6.08
	Qwen3-30B+Tools	7.83	8.43	8.46	6.79
	Qwen3-235B+Tools	8.11	8.78	8.49	7.05
	o3-mini+Code	7.54	8.19	8.13	6.21
RCA	RCA + Qwen2.5	8.09	8.79	8.67	7.13
NOA	RCA + GPT-4.1	8.17	8.93	8.72	7.01

Table 9: Explanation experiment results on CRT dataset abnormal 10%

		CL	LA	EBM	CB
Traditional	Lasso	7.52	8.27	8.16	6.37
MLs	Catboost	7.49	8.21	8.29	6.21
	Qwen2.5-72B	7.56	8.25	8.32	6.73
LLM-Based	DeepSeek-V3	7.48	8.29	8.30	6.02
Methods	DeepSeek-V3.1	7.46	8.30	8.37	6.27
	GPT-4.1	7.32	8.30	8.11	6.37
	DeepSeek-R1	7.83	8.63	8.24	6.81
	o3-mini	7.25	7.86	7.89	5.75
Reasoning	o4-mini	7.13	7.87	7.84	5.59
LLMs	Qwen3-30B	7.62	8.29	7.90	6.43
	Qwen3-235B	7.89	8.62	8.34	6.97
	GPT-5	7.60	8.21	8.13	6.59
	Qwen2.5-72B+Tools	7.59	8.32	8.32	6.83
	DeepSeek-V3+Tools	7.52	8.31	8.22	6.19
	DeepSeek-V3.1+Tools	7.48	8.30	8.08	6.21
	GPT-4.1+Tools	7.29	8.31	8.08	6.25
LLM-Based	DeepSeek-R1+Tools	7.86	8.54	8.26	6.92
Agents	o3-mini+Tools	7.22	7.92	7.98	5.71
	o4-mini+Tools	7.19	7.94	7.95	5.90
	Qwen3-30B+Tools	7.44	8.14	8.22	6.44
	Qwen3-235B+Tools	7.90	8.68	8.65	6.93
	o3-mini+Code	7.52	8.33	8.0	6.84
RCA	RCA + Qwen2.5	7.95	8.75	8.53	6.88
NCA	RCA + GPT-4.1	8.03	8.84	8.69	6.96

Table 10: Complete explanation experiment results in ablation study

			CRT			Diabetes				Heart Disease			
		CL	LA	EBM	СВ	CL	LA	EBM	СВ	CL	LA	EBM	СВ
	original	8.24	8.89	8.47	7.61	8.13	8.57	8.74	6.43	7.62	8.47	8.94	6.18
Qwen2.5-72B	w/o distribution	7.70	8.35	8.10	6.29	7.93	8.47	8.27	6.36	7.32	7.87	8.21	5.74
	w/o reflection	7.54	8.21	7.83	6.24	7.93	8.49	8.11	6.33	7.54	8.23	8.46	5.92
	w/o check	7.75	8.32	8.13	6.45	7.77	8.19	8.52	6.18	7.53	8.16	8.30	5.89
	original	8.16	8.59	8.87	7.62	8.03	8.38	8.63	6.94	7.74	8.53	8.79	6.42
GPT-4.1	w/o distribution	7.45	8.24	7.97	6.19	7.18	7.72	8.24	6.85	7.12	7.93	7.51	5.61
GF 1-4.1	w/o reflection	7.57	8.54	8.06	6.86	7.34	7.94	8.42	6.48	7.13	7.86	7.89	5.76
	w/o check	7.40	8.24	7.83	6.37	7.73	8.06	8.59	6.52	6.72	7.43	6.70	5.10